

Table of contents

- Network architecture
- Node types
- Network discovery
- Full node
- Lightweight node

Network architecture

- Bitcoin network
 - Peer-to-Peer (P2P) network
 - No server/connection hierarchy/centralized node
 - All peers provide and consume services at same
 - Flat mesh topology
 - Originally, Bitcoin network = A collection of nodes running Bitcoin P2P protocol

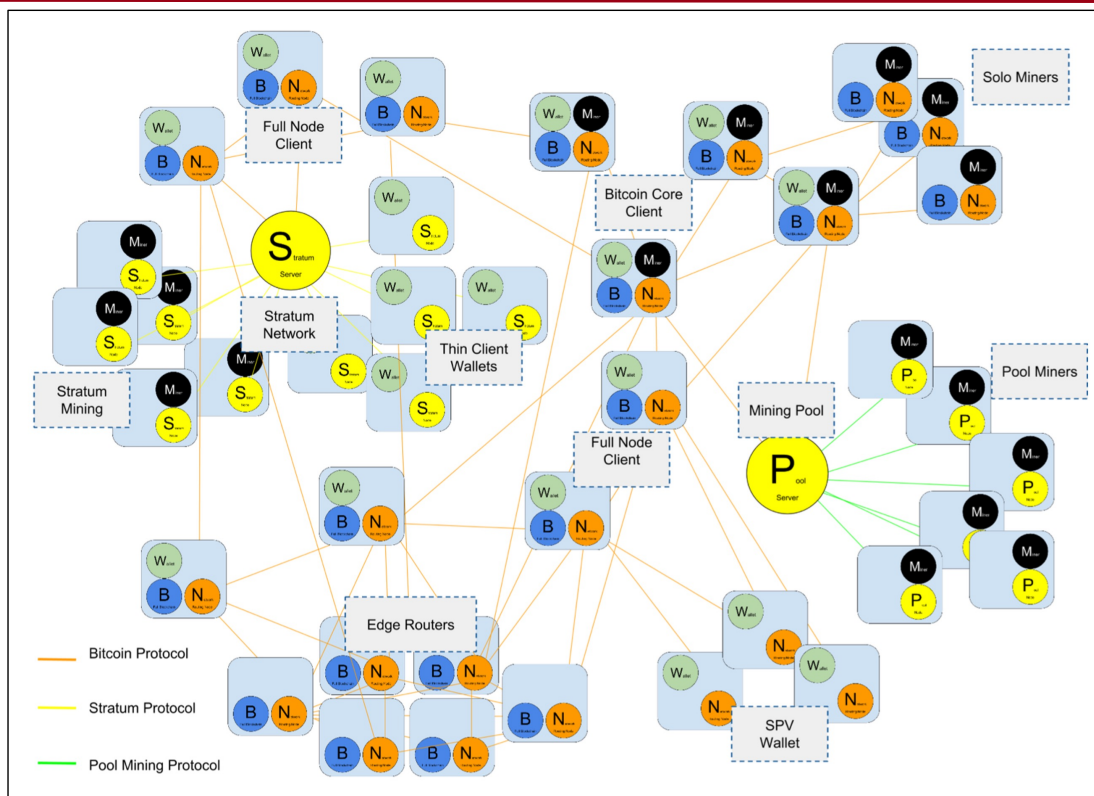
Network architecture

- Bitcoin network
 - Additional protocols (e.g., Stratum) came to support lightweight/mobile wallets
 - Enabled by gateway routing servers
 - Bridge pool miners to Bitcoin P2P
 - Extended Bitcoin Network = Bitcoin P2P protocol + pool mining protocols + Stratum protocol

Node types

Node	Functions	Notation
Reference client (Bitcoin Core or Satoshi Client)	W+M+B+N	W = Wallet M = Miner B = Full blockchain N = Network/routing node P/S = Pool_mining/Stratum
Full node client	W+B+N	
Full blockchain nodes	B+N	
Solo miner	M+B+N	
Lightweight (SPV) wallet	W+N	
Pool protocol servers	P/S server (gateway routers)	
Mining nodes	M+P/S node	
Lightweight (SPV) Stratum wallet	W+S node	

Node types



5

Network discovery

- A booting up node must discover & connect with at least one existing Bitcoin node on network
 - Geographic location of nodes is irrelevant
 - Existing nodes are selected at random
- Establish a TCP connection to a known peer
 - On port 8333 (Bitcoin "known" port) or other provided port
- After connecting, start a "handshake" via *version* message

6

Network discovery

- The *version* message includes
 - *PROTOCOL_VERSION*, constant defining protocol version of initiator
 - *nLocalServices*, list of services supported by initiator, *NODE_NETWORK* for now
 - *nTime*, current time
 - *addrYou*, IP address of remote node as seen by initiator
 - *addrMe*, IP address of initiator, as discovered by initiator
 - *subver*, subversion showing type of software running on initiator
 - *BestHeight*, block height of initiator's blockchain
- Remote node responds with *verack* and establishes a connection
 - Optionally, sends its own *version* message

Network discovery

- How to find peers?
 - No special nodes!
 - Seed nodes
 - Long running stable nodes are listed in clients
 - Not necessary to connect to seed nodes
 - Use them to discover other nodes
 - Give an IP address to booting node

Network discovery


- How to find peers?
 - After one or more connections are established
 - New node sends *addr* message (own IP address) to neighbors
 - Neighbors gossip *addr* message to their neighbors
 - Announcement, visibility
 - New node can also send *getaddr* message to neighbors
 - To get a list of IP addresses of other peers

Full node


- Full node
 - Maintains a complete and up-to-date copy of blockchain with all Tx
 - Independently builds and verifies each block
 - Starting from the genesis block
 - Up to the latest known block in network
 - Keep on building

Full node

- Full node
 - GB of persistent storage required



- Originally, each node was a full node
 - New forms of clients have emerged



11

Lightweight node

- Simplified Payment Verification (SPV)
 - Many devices (e.g., smartphones) are space-/power-constrained
 - Can't store full blockchain
 - SPV enables such devices to work without storing full blockchain
 - Such clients are called SPV/lightweight clients
 - Technique especially used by wallets
 - Different verifications approach
 - Relies on peers for relevant info

12

Lightweight node

- Simplified Payment Verification (SPV)
 - Verifies Tx by reference to their "depth"
 - Instead of their "height" as done by full node
 - When checking a Tx in block #300,000
 - Full node links all 300,000 blocks down to genesis block
 - SPV client verifies it by establishing its depth, say, till block #300,006
 - Since 6 further blocks are built over it (Tx is 6 blocks deep)
 - It isn't a double spend
 - SPV clients
 - Use merkle/authentication path
 - Don't have all Tx/UTXO record
 - Don't download full blocks, just block headers

13

Lightweight node

- Simplified Payment Verification (SPV)
 - SPV client is interested in incoming Tx to addresses in its wallet
 - It establishes a bloom filter on its connections to peers
 - Limit received Tx only to addresses of interest
 - When a peer sees a Tx matching bloom filter
 - Sends a *merkleblock* message, which includes
 - Block header (only header info not entire block)
 - Merkle path of Tx of interest to merkle root in that block
 - As a proof of Tx recording in blockchain, SPV node uses
 - Merkle path, proof of inclusion of Tx in block
 - Block header, linking the block to blockchain

14

Lightweight node

- Simplified Payment Verification (SPV)
 - SPV node consumes <1 KB of data for block header + merkle path
 - Thousand times less than a full block size (=1 MB)
 - SPV node can
 - Definitely prove that Tx exists
 - But, can't verify that a double spend of same UTXO doesn't exist
 - SPV are vulnerable to such double spending attacks
 - To defend, an SPV node connects randomly to several nodes

15

Lightweight node

- Simplified Payment Verification (SPV)
 - SPV nodes retrieve specific/selective Tx
 - Privacy risk
 - Bloom filters
 - Receive a subset of Tx
 - Without revealing precisely addresses of interest
 - Bloom filters
 - Probabilistic data structure
 - An element either is definitely not in the set or may be in the set
 - Correct about actuals/positives (TP=1, FN=0)
 - May be incorrect about fakes/negatives (FP exists)

16