

Table of contents

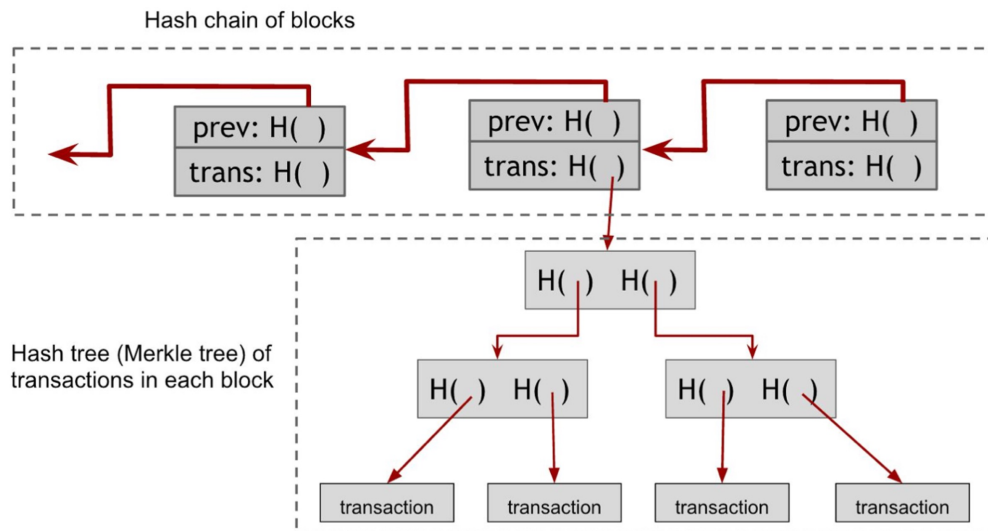
- Bitcoin miners
- Mining hardware
- Mining pool
- Changing protocol

Bitcoin miners

- The task of Bitcoin miners
 - Maintain blockchain
 - Validate Tx
 - Listen for new Tx
 - Assemble a candidate block
 - Find a valid nonce for the candidate block
 - Broadcast
 - Hope block is accepted
 - Profit
 - Validate blocks from the network
 - Listen for new blocks
 - Validate these blocks

Bitcoin miners


- Bitcoin block
 - Bitcoin blockchain contains two different hash structures
 - Hash chain of blocks
 - Merkle tree of Tx within blocks



3

Bitcoin miners

- Bitcoin block
 - Block header is a summary (metadata) of the data in the block



4

Bitcoin miners

- Bitcoin block

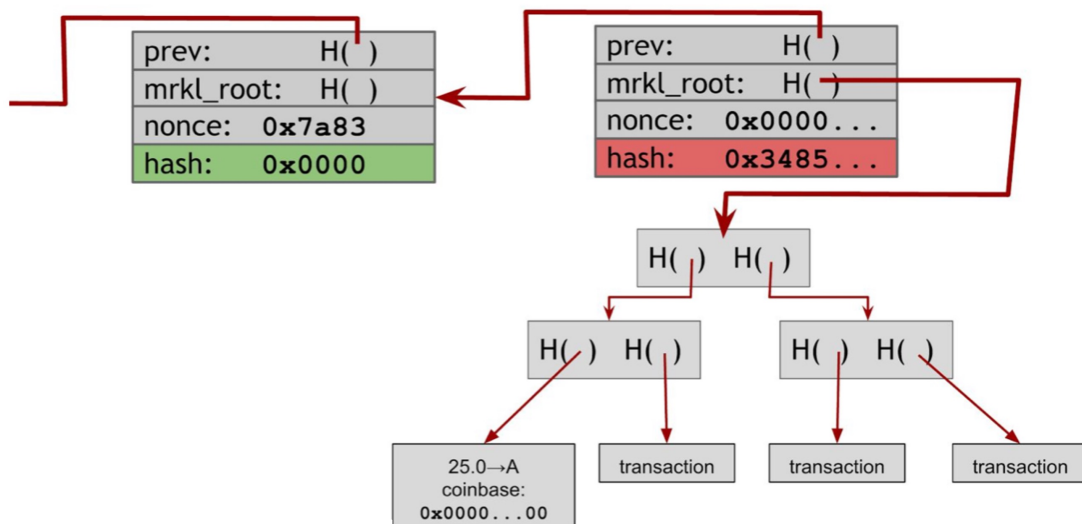
Field	Description
Version	Bitcoin version/rules-set
Previous Block ID	Hash of previous block, where chain is extending
Merkle Root	All Tx in this block hashed (single-line summary of all Tx) together
Time	Unix time, when a miner is trying to mine (hashing) this block the
Bits	A shortened version of target
Nonce	Miners change this 4 bytes (32 bits; 2^{32} possibilities starting from 0) value to get a hash of block header (block hash) below target

Bitcoin miners

- Finding a valid block
 - Miners
 - Identify latest (previous) block
 - Compute the merkle root
 - Pick a set of pending Tx
 - Take care of max size of block
 - Validate these Tx
 - Compile a coinbase Tx
 - Compute target
 - Pick a nonce value
 - Hash the block
 - Check if hash satisfies target
 - If yes, broadcast
 - Otherwise, try again with a new nonce value

Bitcoin miners

- Finding a valid block



7

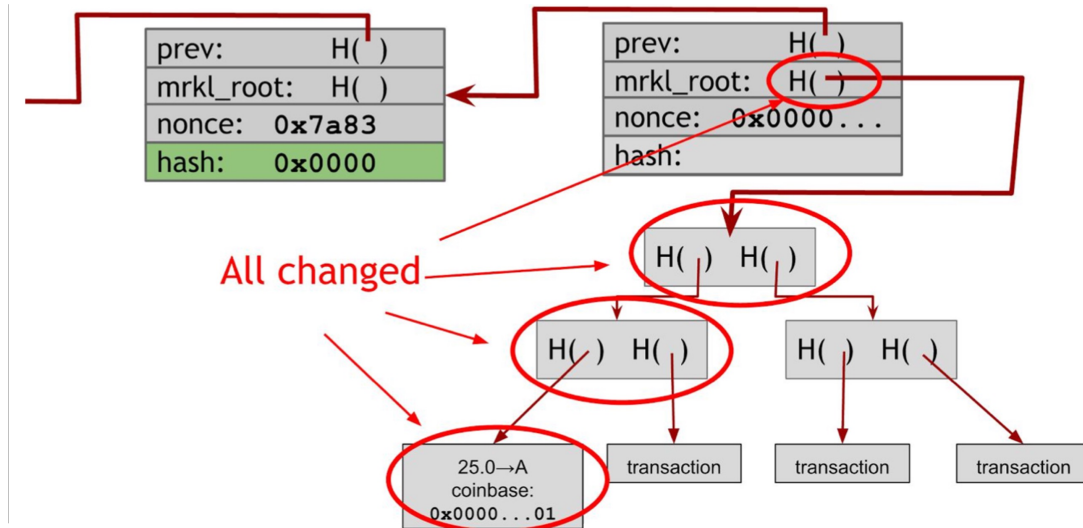
Bitcoin miners

- Finding a valid block
 - What if we exhaust nonce values?
 - Changeable elements in block header
 - Nonce
 - We just exhausted nonce values
 - Time (rollnTime in getwork)
 - Limited scalability
 - Merkle root

8

Bitcoin miners

- Change merkle root
 - Change Tx, costly
 - Change coinbase Tx



9

Bitcoin miners

- Change coinbase Tx
 - In regular Tx
 - Take scriptPubKey from previous Tx and scriptSig from current Tx
 - Coinbase Tx has single blank input
 - No reference to previous output with scriptPubKey
 - scriptSig in current coinbase Tx can be arbitrary bytes
 - Of course, avoiding caveats (*OP_CHECKSIG*)
 - extraNonce solution
 - extraNonce isn't part of the protocol
 - There is no extraNonce field in blocks or Tx

10

Bitcoin miners

- Change coinbase Tx
 - coinBase1/2 represent input/output (overlapping)
 - Change extraNonce (extraNonce1 || extraNonce2)
 - extraNonce1 is unique (pseudo-random) in Stratum (we'll see it later)
 - Only option is to change 4 bytes extraNonce2

version	01000000
input count	01
previous hash	00000000000000000000000000000000 00000000000000000000000000000000
index	fffffffe
scriptlen	60
script	0363600a062f503253482f04358b0553 784404f253000017e46522cfabe6d6d 690682f080c0df0c87cbc7ea4f7f1b5 0005bd0ac3751fc997d9d6971328de 04000000000000004861707079204e59 2120596f7572732047486173682e494f
sequence	00000000
output count	01
value	cb81319500000000
scriptlen	19
script	76a91480ad90d403581fa3bf40686a91 b2d9d4125db6c188ac
lock time	00000000

11

Bitcoin miners

- Change coinbase Tx
 - Overall increment/change extraNonce
 - Coinbase Tx = coinBase1 || extraNonce1 || extraNonce2 || coinBase2
 - Compute $H(\text{coinbase Tx})$
 - Rebuild merkle root
 - Reset nonce, recompute hash

12

Bitcoin miners

- Finding a valid block
 - Is everyone solving the same puzzle?
 - No!
 - Miners choose Tx and build candidate block independently
 - What if different miners choose identical Tx?
 - Blocks would still differ
 - Miners specify their own address in coinbase Tx
 - What if they share the address?
 - Pooled mining
 - Try different ranges of nonce, avoid duplicate work

13

Bitcoin miners

- When you create and broadcast a Tx
 - It goes to a node
 - Node stores/queues it in its mempool
 - Broadcast further
 - Tx reside in mempool until miner picks it
 - Prioritize tx based on fees/age
 - More fee, more chances to be picked
- What if miner never picks it?
 - May be you paid a very low/no fee
 - Miner is least interested to pick it
 - Can I increase fee?
 - Replace By Fee (RBF)

14

Bitcoin miners

- RBF
 - A node policy
 - Allows an unconfirmed Tx in a mempool to be replaced with a different Tx
 - Spends at least one of same inputs
 - Pays a higher Tx fee

Bitcoin miners

- RBF
 - Full RBF
 - Unconditionally allows a Tx to replace older ones as long as it pays a sufficient fee
 - Opt-in RBF
 - Only allows replacement when Tx being replaced have explicitly signaled to allow replacement
 - Via nSequence/sequence field
 - Both full and opt-in RBF may lead to double spend

Bitcoin miners

- RBF
 - First-seen-safe RBF
 - Only allows replacement if replacement/new Tx keeps same outputs as Tx being replaced
 - Delayed RBF
 - Allows Tx to be replaced unconditionally, but only after a given number of blocks have been mined since Tx being replaced were first seen by node

Bitcoin miners

- RBF
 - nSequence
 - 4 bytes field in an Tx input
 - Usually only shown in rawTx
 - Send a Tx with a lock_time & sequence number 0
 - Tx is then not considered "final" by network
 - Can't include in a block until lock_time
 - Before lock_time expires
 - Replace Tx with as many new versions as you want
 - Newer versions have higher sequence numbers
 - To lock Tx permanently
 - Set sequence number to 0xFFFFFFFF
 - Tx is considered final, even if lock_time has not reached

Bitcoin miners

- Orphan Tx
 - When Tx are transmitted across network
 - May arrive to nodes in different order
 - A child might arrive before parent
 - Upon seeing a Tx (child)
 - That references a Tx (parent) that is not yet known
 - Nodes don't reject child
 - Put it in temporary orphan Tx pool
 - Where orphan Tx wait for parent's arrival
 - When parent arrives
 - Child Tx is retrieved, revalidated recursively, inserted in mempool

19

Bitcoin miners

- Orphan block
 - When a fork occurs
 - A block may not be included in main chain
 - Block becomes orphaned
 - All Tx in it simply go back in node's queue/pool
 - Revalidated, if their input has been spent
 - If so, evicted
 - Wait to be added to next block

20

Mining hardware

- Computing SHA-256 hash function is the core of miners' tasks
- CPU mining
 - First generation mining was done on general purpose CPUs
 - Simply compute SHA-256 in software, check the result
 - Speed?
 - ~10-20 MH/s
 - Mining on a regular CPU at current difficulty level is impossible

21

Mining hardware

- GPU mining
 - Designed for processing videos/graphics
 - Large number of floating point units
 - Not used at all in SHA-256
 - Poor cooling characteristics (especially, in farms)
 - High electricity computation
 - Speed?
 - ~200-300 MH/s
 - Ok, but not sufficient

22

Mining hardware

- FPGA mining
 - Better performance over GPUs
 - "Bit fiddling" operations is trivial to specify
 - Use all transistors on the card
 - Cooling is easier
 - Speed?
 - ~1 GH/s
 - Better, but still not sufficient
 - Cost-per-performance was marginally improved over GPUs
 - Short-lived phenomenon

23

Mining hardware

- ASIC mining
 - Application-Specific Integrated Circuits
 - Optimized for the sole purpose of mining
 - Lifetime of ASICs was quite short initially
 - Due to rapidly increasing network hash rate
 - Now ASIC equipments have longer life time
 - After growth rate of hash power has stabilized

24

Mining hardware

- Economics of mining haven't been favorable to small miner
 - Mining cost = H/w cost + operation cost (electricity, cooling, etc.)
- Today, professional mining
 - Mining farms
 - Mining pools
- Future, find alternatives

Mining pool

- A group of miners form a pool
- Attempt to mine a block
- Use a designated coinbase recipient
 - Recipient is called pool manager
- No matter who actually finds nonce
 - Pool manager will receive reward


Mining pool

- Pool manager takes that revenue
 - Distribute it among participants
 - Based on amount of work done by participants
 - Takes a cut
- How does pool manager know if miners are actually working?
 - If yes, how to estimate amount of work?
 - Miners can cheat
 - Show proof?

27

Mining pool

- Mining shares
 - Miners probabilistically prove amount of work done
 - Consider a target begins with 30 zeros
 - Block's valid hash must have ≥ 30 zeros
 - Miners find many hashes, some are close to target
 - Miners show these nearly valid hashes as proof
 - Threshold can be set, e.g., ≥ 20 zeros
 - More work, more profit
 - No bonus for finding right nonce



28

Mining pool

- Mining shares
 - Pay-per-share model
 - Pool manager pays a flat fee for shares above threshold
 - Best for miners
 - Pool manager absorbs the risk
 - Must pay miners even if right nonce is not found
 - Proportional model
 - Payment depends on whether or not pool actually finds right nonce
 - If found, reward is distributed based on contributions
 - Otherwise, no/low reward
 - Miners do pool hopping to optimize revenue

29

Mining pool

- Are mining pools good?
 - Pros
 - Easier for smaller miners to get involved
 - Easier to upgrade network
 - Central pool manager can force miners to upgrade
 - Cons
 - Issue of centralization
 - 51% mining pools
 - All are miners, less full nodes

30

Mining pool

- Mining pool protocols
 - Three major protocols
 - getwork
 - getblocktemplate (gbt)
 - Stratum

31

Mining pool

- Mining pool protocols
 - getwork
 - Original mining method
 - HTTP-based requests and responses
 - JSON Remote Procedure Calls (RPC) to request a block
 - Server gives only block headers to client
 - No Tx or way to modify block (except nonce)
 - Client is limited to try all nonce values
 - If exhausted, get more work from server
 - High bandwidth usage for modern mining h/w
 - Worst approach in today's scenario
 - rollnTime extension
 - Allows a client to generate own work by modifying timestamp on block header
 - Limited scalability

32

Mining pool

- Mining pool protocols
 - getblocktemplate (gbt)
 - HTTP-based
 - JSON RPC for communication between client and poolserver
 - Server gives templates to client
 - Client uses it to generate own work
 - Full block data is given
 - Allows the client to modify the block
 - Miner can choose which Tx to include in block
 - Very high bandwidth usage
 - Well documented, detailed specifications

33

Mining pool

- Mining pool protocols
 - Stratum
 - Not HTTP-based
 - Direct TCP connections
 - JSON RPC for communication between client and poolserver
 - Uses parts of gbt's message specification
 - Server gives templates to client
 - Client uses it to generate own work
 - Block headers and only Tx hashes are given
 - Change extraNonce
 - Least bandwidth usage
 - Very fast, efficient switching to new work
 - No real specification

34

Mining pool

- Mining pool protocols

- Stratum

```
{
    "result": [
        {
            "mining.notify",
            "ae6812eb4cd7735a302a8a9dd95cf71f"
        },
        {
            "60021014", extraNonce1
        },
        {
            "extraNonce2 size (bytes)"
        }
    ],
    "id": 1,
    "error": null
}

{
    "params": [
        {
            "Difficulty"
        },
        {
            "id": null,
            "method": "mining.set_difficulty"
        }
    ]
}

{
    "params": [
        {
            "JobID"
        },
        {
            "Previous block hash"
        },
        {
            "Merkle branches"
        }
    ],
    "id": null,
    "method": "mining.notify"
}
```

35

Changing protocol

- How to introduce new features in Bitcoin protocol?
 - Just release an update and ask all nodes to upgrade?
 - Not really!
 - Not every node would upgrade
 - Some nodes would fail to get or to get it in time
 - Two types of changes
 - One causes a hard fork
 - Other causes a soft fork

36

Changing protocol

- Hard fork
 - Introduces new features that were previously considered invalid
 - New nodes would accept blocks as valid that old nodes would reject
 - After majority of nodes have upgrade
 - Longest branch will contain blocks that are invalid for old nodes
 - So, old nodes will consider their (shorter) branch to be longest valid branch
 - Such changes cause a hard fork
 - Blockchain splits
 - Branches will never join together again

37

Changing protocol

- Hard fork
 - E.g., fixing the bug in *OP_CHECKMULTISIG* would need a hard fork
 - What can we add with a hard fork
 - Adding new opcodes to protocol
 - Changing limits on block/Tx size

38

Changing protocol

- Soft fork
 - Adding features that make validation rules stricter
 - Old node would accept all blocks while new nodes would reject some
 - After majority of nodes have upgraded
 - New nodes will enforce new, tighter rules
 - Old miners might mine invalid blocks
 - Include some Tx invalid under new rules
 - But, they will figure out soon after their blocks are rejected
 - Without understanding reason
 - Ask their operator to check/upgrade
 - Such changes cause a soft fork
 - There will be many small, temporary forks

39

Changing protocol

- Soft fork
 - E.g., introduction of P2SH scripts
 - Old nodes would still verify a valid P2SH correctly
 - Hash the data value, check if hash matches value in scriptPubKey
 - OP_EQUAL pushes 1 on top of stack
 - Old nodes don't know (now required) step of running that value itself
 - What can we add with a soft fork
 - New cryptographic schemes
 - Some extra metadata in coinbase parameter
 - Some specific formatting to a field
 - 2016 “The DAO” hack
 - Soft/hard fork proposals

40